

Processes

processes practice for CSC420 Operating Systems Spring 2022 Lyon College

README

- This file accompanies lectures on the shell and `bash(1)`. To gain practice, you should type along in your own Org-mode file. You have to have Emacs and my `.emacs` file installed on your PC or the Pi you're working with.
- This section is based on chapter 10 of Shotts, *The Linux Command Line* (2e), NoStarch Press (2019).
- To make this easier, use the auto expansion (`<s`). This will only work if you have my `.emacs` file ([from GDrive](#)) installed.
- Add the following two lines at the top of your file, and activate each line with `C-c C-c` (this is confirmed in the echo area as `Local setup has been refreshed`):

```
#+PROPERTY: header-args: bash :results output
```

- Remember that `C-M-\` inside a code block indents syntactically (on Windows, this may only work if you have a marked region - set the mark with `C-SPC`).

Hide Org-mode emphasis markers

- []

Add the string `"(setq-default org-hide-emphasis-markers t)"` to your `$HOME/.emacs` file so that you don't see the emphasis markers in Org-mode. Then show the last line of the file to make sure it worked.

Tip: you need the commands `echo`, `~>>`, `tail`.

```
echo "(setq-default org-hide-emphasis-markers t)" >> $HOME/.emacs
tail -1 $HOME/.emacs
```

```
(setq-default org-hide-emphasis-markers t)
```

Overview

- Modern operating systems are *multitasking*, which means they create an illusion of doing more than one thing at once.
- They do this by rapidly switching from one executing program to another.
- The **kernel** manages this through clever **process management**, which really is clever **memory management**.
- This is illustrated in the figure [1](#). The simple program is all over the computer's memory.



Figure 1: Memory Layout of a C Program (Source: Silberschatz et al)

- As always, let's focus on stuff we can do ourselves. This includes a bunch of new shell commands:

COMMAND	MEANING
ps	Report a snapshot of current processes
top	Display tasks
jobs	List active jobs
bg	Place a job in the background
fg	Place a job in the foreground
kill	Send a kill signal to process
killall	Kill processes by name
shutdown	Shut down or reboot the system

How a process works

- When the OS starts up, the **kernel** launches the **init** program, which in turn runs a series of shell scripts (in `/etc`) that start all the system services.
- [] Check `/etc` out now - you find e.g. the directory `/etc/cups`, which contains scripts for the Common UNIX Printing System (CUPS).
- Many of the services are *daemon* programs - they just sit in the background and do their thing without a user interface (UI).
- `init` itself is a daemon, also called `systemd`. The shell program `systemctl` allows indirect access to all services.
- []

Grab a daemon!

In the code block [1](#),

1. run the command `systemctl status`,
2. tee its output to a text file `systemctl.txt`
3. `grep` for the login daemon program `logind`

```
systemctl status | tee systemctl.txt | grep logind
```

```
|      └─1618 grep logind
|      └─systemd-logind.service
|      └─374 /lib/systemd/systemd-logind
```

- If a program (like `init`) can launch other programs, it's a *parent process* producing a *child process*.
- **How does the kernel maintain control?** By assigning a *process ID* (PID) to every process.
- Processes are assigned in ascending order beginning with `init`, which has PID 1.
- The **kernel** also tracks process memory and readiness to resume execution. Like files, processes have owners and userIDs.

Viewing processes statically

- The `ps` program has a lot of options (check `ps(1)`)
- []

Run `ps` without options.

```
ps
```

```
PID TTY          TIME CMD
853 ?            00:00:00 systemd
854 ?            00:00:00 (sd-pam)
881 ?            00:00:00 pipewire
882 ?            00:00:00 pulseaudio
931 ?            00:00:00 dbus-daemon
932 ?            00:00:00 pipewire-media-
1016 ?           00:00:00 ssh-agent
1096 ?           00:00:00 gvfsd
1101 ?           00:00:00 gvfsd-fuse
1127 ?           00:00:00 ssh-agent
1179 ?           00:00:00 menu-cached
1183 ?           00:00:00 gvfs-udisks2-vo
1201 ?           00:00:00 gvfs-afc-volume
1206 ?           00:00:00 gvfs-gphoto2-vo
1210 ?           00:00:00 gvfs-mtp-volume
1214 ?           00:00:00 gvfs-goa-volume
1228 ?           00:00:00 gvfsd-trash
1619 ?           00:00:00 bash
1620 ?           00:00:00 ps
```

- The result is confusing because you're inside another program now.
- []

Open a shell (in Emacs with `M-x shell` or a terminal) and type `ps`. You should see something like this:

```
PID TTY          TIME CMD
12254 pts/1      00:00:00 bash
12257 pts/1      00:00:00 ps
```

- **What this means:**
 - You see two PID - the shell program and the `ps` program
 - TTY ("teletype") is the *controlling terminal* for the process
 - TIME is the amount of CPU time consumed by the process
- []

Run `ps` again, this time add the option `x`

```
ps x
```

```

PID TTY      STAT   TIME COMMAND
853 ?        Ss      0:00 /lib/systemd/systemd --user
854 ?        S        0:00 (sd-pam)
881 ?        S<sl    0:00 /usr/bin/pipewire
882 ?        S<sl    0:00 /usr/bin/pulseaudio --daemonize=no --log-target=journal
915 tty1     S        0:00 -bash
931 ?        Ss      0:00 /usr/bin/dbus-daemon --session --address=systemd: --nofo
932 ?        S<l     0:00 /usr/bin/pipewire-media-session
939 tty1     S+       0:00 /bin/sh /usr/bin/startx
961 tty1     S+       0:00 xinit /etc/X11/xinit/xinitrc -- /etc/X11/xinit/xserverrc
962 tty1     Sl       0:24 /usr/lib/xorg/Xorg -nolisten tcp :0 vt1 -keeptty -auth /
967 tty1     Sl       0:00 /usr/bin/lxsession -s LXDE-pi -e LXDE
1016 ?       Ss      0:00 /usr/bin/ssh-agent x-session-manager
1096 ?       Ssl     0:00 /usr/libexec/gvfsd
1101 ?       Sl       0:00 /usr/libexec/gvfsd-fuse /run/user/1000/gvfs -f
1113 tty1     Sl       0:19 mutter
1117 tty1     Sl       0:00 lxdm
1118 tty1     Sl       0:03 lxpanel --profile LXDE-pi
1119 tty1     Sl       0:01 pcmanfm --desktop --profile LXDE-pi
1127 ?       Ss      0:00 /usr/bin/ssh-agent -s
1128 tty1     S        0:00 /usr/bin/python3 /usr/share/system-config-printer/applet
1179 ?       Sl       0:00 /usr/lib/menu-cache/menu-cached /run/user/1000/menu-cach
1183 ?       Ssl     0:00 /usr/libexec/gvfs-udisks2-volume-monitor
1201 ?       Ssl     0:00 /usr/libexec/gvfs-afc-volume-monitor
1206 ?       Ssl     0:00 /usr/libexec/gvfs-gphoto2-volume-monitor
1210 ?       Ssl     0:00 /usr/libexec/gvfs-mtp-volume-monitor
1214 ?       Ssl     0:00 /usr/libexec/gvfs-goa-volume-monitor
1228 ?       Sl       0:00 /usr/libexec/gvfsd-trash --spawner :1.7 /org/gtk/gvfs/ex
1439 tty1     Sl       0:38 /usr/bin/emacs
1508 pts/0    Ss+     0:00 /bin/bash --noediting -i
1621 ?       Ss      0:00 bash
1622 ?       R        0:00 ps x

```

- `ps x` (no dash!) shows all processes regardless of what terminal they are controlled by. `?` indicates no terminal (like daemons).
- []

How many processes that you own that have no terminal?

```
ps x | grep ? | wc -l
```

```
21
```

- []

List only the first 5 lines of the `ps x` listing.

```
ps x | head -5
```

```

PID TTY      STAT   TIME COMMAND
853 ?        Ss      0:00 /lib/systemd/systemd --user

```

```

854 ?      S      0:00 (sd-pam)
881 ?      S<sl   0:00 /usr/bin/pipewire
882 ?      S<sl   0:00 /usr/bin/pulseaudio --daemonize=no --log-target=journal

```

- The column STAT reveals the current status of the process, see table [1](#).

STATE	MEANING
R	Running or ready to run
S	Sleeping, waiting for an event (e.g. keystroke)
D	Uninterruptible sleep, waiting for I/O (e.g. disk)
T	Stopped, received instruction to stop
Z	Zombie child process, abandoned by parent
<	High priority (not <i>nice</i> - more CPU time)
N	Low priority (<i>nice</i>) - served once < are done

There may be more characters denoting exotic process characteristics (see `ps(1)`). E.g. *s* is a *session leader*, *+* is a *foreground* process, and *l* is multi-threaded.

- []

You get even more information with the option `aux`. Redirect the output of `ps aux` to a file `psaux.txt`, and print only the first 5 lines.

```
ps aux | tee psaux.txt | head -5
```

```

USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root         1   0.3   0.2 33708   8680 ?        Ss   22:45   0:03 /sbin/init splash
root         2   0.0   0.0      0      0 ?        S    22:45   0:00 [kthreadd]
root         3   0.0   0.0      0      0 ?        I<   22:45   0:00 [rcu_gp]
root         4   0.0   0.0      0      0 ?        I<   22:45   0:00 [rcu_par_gp]

```

- You should see PID 1, the `init` program. The `splash` options means that you can see a splash screen during boot.
- Table [1](#) shows some header definitions

HEADER	MEANING
USER	User ID - this is the process owner
%CPU	CPU usage in percent
%MEM	Memory usage in percent
VSZ	Virtual memory size
RSS	Resident set size - RAM use in kB

HEADER	MEANING
START	Process starting time and date

- []

Why is the CPU usage of `init` zero, while the Memory usage is non-zero? How much RAM does the program actually use?

ANSWER: The `init` program only runs during the booting process, but as part of the **kernel** it is loaded into the central memory. It occupies 8MB.

Viewing processes dynamically

- `ps` provides a snapshot, but `top` provides a real-time view.
- [] Open a terminal (in or outside of Emacs) and run `top`. You can stop the command with `C-c` or `q`.
- `top` refreshes every three seconds and shows the top system processes. It includes a summary at the top and a table sorted by CPU activity at the bottom.



Figure 2: Top view

- The system summary contains a lot of good stuff. Table [1](#) gives a rundown.

ROW	FIELD	MEANING
1	<code>top</code>	Program name
	21:52:54	Current time of day
	up 2 days 9:49	<i>uptime</i> since last boot
	1 user	No. of users logged in
	load average	No. of processes waiting to run Values < 1.0 means not busy
2	Tasks:	No. of processes and their states total, running, sleeping, stopped
3	Cpu(s):	Activities that the CPU performs: us: user processes (not kernel) sy: system processes (kernel) ni: nice (low prio) processes id: idle processes wa: waiting for I/O

ROW	FIELD	MEANING
4	Mem:	Physical RAM used
5	Swap:	Swap space (virtual memory) used

- top accepts some keyboard commands like h (help) and q (quit).
- top is better than any graphical application (e.g. the Task Manager that you have on your Pi) - it is faster and consumes far less resources.

Controlling processes

Interrupting a process

- As a guinea pig program, we use emacs.
- [] Open a terminal (inside Emacs after splitting the screen with C-x 2 or outside of Emacs), and enter emacs at the prompt. A new Emacs editor window appears. Notice that the terminal prompt does not return.
- [] Close the new Emacs editor manually by clicking on the x in the upper right corner. The prompt in the Shell returns.
- [] Enter emacs again in the shell, and interrupt it with CTRL-C (outside of Emacs, or with C-c C-c on the Emacs *shell*).
- Many programs can be interrupted this way by sending an **interrupt** signal to the **kernel**.

Putting a process in the background

- The terminal has a *foreground* and a *background*. To launch a program so that it is immediately placed into the background, follow it with an ampersand & character
- [] Start Emacs from the shell in the background. An Emacs window should open. Look at the terminal.
- The message that appeared is part of shell *job control*. It means that we have started job number 1 with the PID 13899. If you check the process table with ps, you should see the process

```
[1] 13899
```

- []

grep the emacs process from the process table using the PID.

```
13928 pts/1    00:00:04 emacs
```

- []

The jobs command lists the jobs that were launched from our terminal. Try it. You should see something like this:

```
[1]+  Running                  emacs &
```

Returning a process to the foreground

- A process in the background is immune from keyboard input - you cannot interrupt it with CTRL-C. To return it to the foreground, use the `fg` command.
- [] On the shell where you started it, return the process to the foreground with the command `fg %1`. The 1 is the jobspec.
- [] Kill the Emacs process with `C-c C-c` or CTRL-C on the shell where you started it.

Stopping or pausing a process

- []

Start an emacs process in a terminal (NOT in an Emacs shell) - it's now in the foreground. If you press CTRL-Z in the shell, the process is stopped.

```
pi@raspberrypi:~ $ emacs
^Z
[1]+  Stopped                  emacs
pi@raspberrypi:~ $
```

- [] To bring the process back, you can either bring it into the foreground with `fg %1`, or resume the process in the background with `bg %`. Try both.
- **Why would you launch a graphical program from the shell?**
 - The program may not be listed in the GUI
 - You see error messages that otherwise are invisible
 - Some graphical programs have useful command line options

Killing a process

- []

`kill` is used to terminate processes using the PID. Start Emacs from the shell *in the background* (inside or outside of emacs), and then kill it with `kill PID`.

Tip: you get the PID with `ps`, or right after executing the background command.

- `kill` does actually not "kill" the process, it sends it a signal. We have already used some of these signals:

SIGNAL	MEANING
INT	CTRL-C - interrupt process
TSTP	CTRL-Z - terminal stop
HUP	Hang up (used by daemons)
KILL	Kill without cleanup
TERM	Terminate with <code>kill</code>
STOP	Stop without delay

- Some of these signals are sent to the target program (identified by PID) while others are sent straight to the kernel.

More process commands

Some fun commands to play with and explore. We already looked at `ps tree`. You may have to install these.

COMMAND	MEANING
<code>ps tree</code>	Process list arranged as tree pattern
<code>vmstat</code>	System usage snapshot
<code>xload</code>	Draws a graph showing system load over time
<code>tlod</code>	Draws graph in terminal

Summary

References

- Silberschatz, Galvin and Gagne (2018). Operating System Concepts - 10th edition, Wiley.

Author: Marcus Birkenkrahe

Created: 2022-04-04 Mon 23:01

[Validate](#)